

IPS9 in R: Looking at Data – Relations (Chapter 2)

Margaret Chien and Nicholas Horton (nhorton@amherst.edu)

July 19, 2018

Introduction and background

These documents are intended to help describe how to undertake analyses introduced as examples in the Ninth Edition of *Introduction to the Practice of Statistics* (2017) by Moore, McCabe, and Craig.

More information about the book can be found [here](#). The data used in these documents can be found under Data Sets in the Student Site. This file as well as the associated R Markdown reproducible analysis source file used to create it can be found at <https://nhorton.people.amherst.edu/ips9/>.

This work leverages initiatives undertaken by Project MOSAIC (<http://www.mosaic-web.org>), an NSF-funded effort to improve the teaching of statistics, calculus, science and computing in the undergraduate curriculum. In particular, we utilize the `mosaic` package, which was written to simplify the use of R for introductory statistics courses. A short summary of the R needed to teach introductory statistics can be found in the mosaic package vignettes (<http://cran.r-project.org/web/packages/mosaic>). A paper describing the mosaic approach was published in the *R Journal*: <https://journal.r-project.org/archive/2017/RJ-2017-024>.

Chapter 2: Looking at Data – Relationships

This file replicates the analyses from 2: Looking at Data – Relationships.

First, load the packages that will be needed for this document:

```
library(mosaic)
library(readr)
```

Section 2.1: Relationships

Section 2.2: Scatterplots

Example 2.8: Laundry detergents

```
Laundry <- read_csv("https://nhorton.people.amherst.edu/ips9/data/chapter02/EG02-008LAUNDRY.csv")
```

```
## Parsed with column specification:
## cols(
##   Price = col_integer(),
##   Rating = col_integer(),
##   Type = col_character()
## )
```

```
# 2.10: Examine the spreadsheet
str(Laundry)
```


```
## Classes 'tbl_df', 'tbl' and 'data.frame':   53 obs. of  3 variables:
## $ Price : int  14 14 14 12 10 8 56 14 18 13 ...
## $ Rating: int  60 58 60 53 57 41 51 29 44 65 ...
## $ Type : chr  "Liquid" "Liquid" "Liquid" "Liquid" ...
## - attr(*, "spec")=List of 2
```

```
## ..$ cols      :List of 3
## .. ..$ Price : list()
## .. .. ..- attr(*, "class")= chr  "collector_integer" "collector"
## .. ..$ Rating: list()
## .. .. ..- attr(*, "class")= chr  "collector_integer" "collector"
## .. ..$ Type   : list()
## .. .. ..- attr(*, "class")= chr  "collector_character" "collector"
## ..$ default: list()
## .. ..- attr(*, "class")= chr  "collector_guess" "collector"
## ..- attr(*, "class")= chr "col_spec"
```

```
favstats(~ Rating, data = Laundry)
```

```
## min Q1 median Q3 max      mean      sd n missing
##  25 42      51 61 85 53.01887 14.25387 53      0
```


```
gf_histogram(~ Rating, data = Laundry, binwidth = 5)
```



```
favstats(~ Price, data = Laundry)
```

```
## min Q1 median Q3 max      mean      sd n missing
##   5 12      14 24 56 17.37736 8.838783 53      0
```

```
gf_histogram(~ Price, data = Laundry, binwidth = 5)
```



Example 2.9: Laundry detergents

```
Laundry <- read_csv("https://nhorton.people.amherst.edu/ips9/data/chapter02/EG02-009LAUNDRY.csv")
```

```
## Parsed with column specification:
## cols(
##   Price = col_integer(),
##   Rating = col_integer(),
##   Type = col_character()
## )
```

Figure 2.1, page 86

```
gf_point(Rating ~ Price, data = Laundry, xlab = "Price per load (cents)")
```





Figure 2.2


```
gf_point(Rating ~ Price, data = Laundry, xlab = "Price per load (cents)") +
  xlim(0, 35)
```

```
## Warning: Removed 1 rows containing missing values (geom_point).
```



```
# 2.12: Make a scatterplot
gf_point(jitter(Rating) ~ jitter(Price), data = Laundry, xlab = "Price per load (cents)") +
  xlim(0, 35)
```

```
## Warning: Removed 1 rows containing missing values (geom_point).
```




We can use `jitter()` to add some noise into the plot to show overlapped points.

```
# 2.13: Change the units
Laundry2 <- Laundry %>%
  mutate(Price = Price/100)
favstats(~ Price, data = Laundry2)
```

```
##   min   Q1 median   Q3   max     mean      sd  n missing
## 0.05 0.12   0.14 0.24 0.56 0.1737736 0.08838783 53      0
```

```
gf_point(Rating ~ Price, data = Laundry2, xlab = "Price per load (dollars)")
```




Example 2.10: Scatterplot with a straight line

```
Laundry <- read_csv("https://nhorton.people.amherst.edu/ips9/data/chapter02/EG02-010LAUND.csv")
```

```
## Parsed with column specification:
## cols(
##   Price = col_integer(),
##   Rating = col_integer(),
##   Type = col_character()
## )
```

```
gf_point(Rating ~ Price, data = Laundry, xlab = "Price per load (cents)") %>%
  gf_lm()
```



Example 2.11: Education spending and population: Benchmarking

```
EduSpending <- read_csv("https://nhorton.people.amherst.edu/ips9/data/chapter02/EG02-011EDSPEND.csv")
```


```
## Parsed with column specification:
## cols(
##   State = col_character(),
##   Spending = col_double(),
##   Population = col_double()
## )
```

```
head(EduSpending)
```

```
## # A tibble: 6 x 3
##   State      Spending Population
##   <chr>      <dbl>      <dbl>
## 1 Alabama      14.9         4.9
## 2 Alaska       3.8         0.7
## 3 Arizona      14.8         6.8
## 4 Arkansas      8.5         3
## 5 California  111.        39.2
## 6 Colorado     14.3         5.4
```

Figure 2.5, page 90

```
gf_point(Spending ~ Population, data = EduSpending) %>%
  gf_lm() %>%
  gf_labs(x = "Population (millions)", y = "Spending ($ billion)")
```




Example 2.12: Calcium retention

```
Calcium <- read_csv("https://nhorton.people.amherst.edu/ips9/data/chapter02/EG02-012CALCIUM.csv")
```

```
## Parsed with column specification:
## cols(
##   CaIntake = col_double(),
##   CaRetention = col_double(),
##   LogRet = col_double()
## )
```

```
# Figure 2.6
gf_point(CaRetention ~ CaIntake, data = Calcium) %>%
  gf_smooth() %>%
  gf_labs(x = "Calcium intake (mg/d)", y = "Calcium retention (mg/d)")
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```




Example 2.13: Calcium retention with logarithms

```
Calcium <- read_csv("https://nhorton.people.amherst.edu/ips9/data/chapter02/EG02-013CALCIUM.csv")
```

```
## Parsed with column specification:
## cols(
##   CaIntake = col_double(),
##   CaRetention = col_double(),
##   LogRet = col_double()
## )
```

```
# Figure 2.7, page 91
gf_point(LogRet ~ CaIntake, data = Calcium) %>%
  gf_lm() %>%
  gf_labs(x = "Calcium intake (mg/d)", y = "Log calcium retention")
```




Example 2.14: Education spending and population with logarithms

```
EduSpending <- read_csv("https://nhorton.people.amherst.edu/ips9/data/chapter02/EG02-014EDSPEND.csv")
```

```
## Parsed with column specification:
## cols(
##   State = col_character(),
##   Spending = col_double(),
##   Population = col_double()
## )
```

Figure 2.8, page 92

```
EduSpending %>%
  mutate(LogPop = log(Population), LogSpend = log(Spending)) %>%
  gf_point(LogSpend ~ LogPop) %>%
  gf_lm() %>%
  gf_labs(x = "Log population", y = "Log spending")
```




Example 2.15: Rating versus price and type of laundry detergent

```
Laundry <- read_csv("https://nhorton.people.amherst.edu/ips9/data/chapter02/EG02-015LAUND.csv")
```

```
## Parsed with column specification:
## cols(
##   Price = col_integer(),
##   Rating = col_integer(),
##   Type = col_character()
## )
```

Figure 2.9, page 93

```
gf_point(Rating ~ Price, color = ~ Type, data = Laundry, xlab = "Price per load (cents)") %>%
  gf_lm()
```



Example 2.16: Laundry rating versus price with a smooth fit


```
Laundry <- read_csv("https://nhorton.people.amherst.edu/ips9/data/chapter02/EG02-016LAUND.csv")
```

```
## Parsed with column specification:
## cols(
##   Price = col_integer(),
##   Rating = col_integer(),
##   Type = col_character()
## )
```

Figure 2.10, page 94-95


```
gf_point(Rating ~ Price, data = Laundry) %>%
  gf_smooth(span = .5)
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



```
gf_point(Rating ~ Price, data = Laundry) %>%
  gf_smooth(span = 5)
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



Example 2.17: A smooth fit for education spending and population with logs

```
EduSpending <- read_csv("https://nhorton.people.amherst.edu/ips9/data/chapter02/EG02-017EDSPEND.csv")
```


```
## Parsed with column specification:
## cols(
##   State = col_character(),
##   Spending = col_double(),
##   Population = col_double()
## )
```

```
# Figure 2.8, page 92
```

```
EduSpending %>%
  mutate(LogPop = log(Population), LogSpend = log(Spending)) %>%
```

```
gf_point(LogSpend ~ LogPop) %>%
gf_smooth() %>%
gf_labs(x = "Log population", y = "Log spending")
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```



Section 2.3: Correlation

Use Your Knowledge: Laundry detergents

```
# page 102
cor(Rating ~ Price, data = Laundry)
```

```
## [1] 0.2109681
```


Section 2.4: Least Squares Regression

Example 2.19: Fidgeting and fat gain

```
Fidgeting <- read_csv("https://nhorton.people.amherst.edu/ips9/data/chapter02/EG02-019FIDGET.csv")
```

```
## Parsed with column specification:
## cols(
##   NEA = col_integer(),
##   Fat = col_double(),
##   Resid = col_double()
## )
```

```
# Figure 2.16, page 108
gf_point(Fat ~ NEA, data = Fidgeting) %>%
gf_labs(x = "Nonexercise activity (calories)", y = "Fat gain (kilograms)")
```




Example 2.20: Regression line for fat gain

```
Fidgeting <- read_csv("https://nhorton.people.amherst.edu/ips9/data/chapter02/EG02-020FIDGET.csv")
```

```
## Parsed with column specification:
## cols(
##   NEA = col_integer(),
##   Fat = col_double(),
##   Resid = col_double()
## )
```

Figure 2.17, page 109

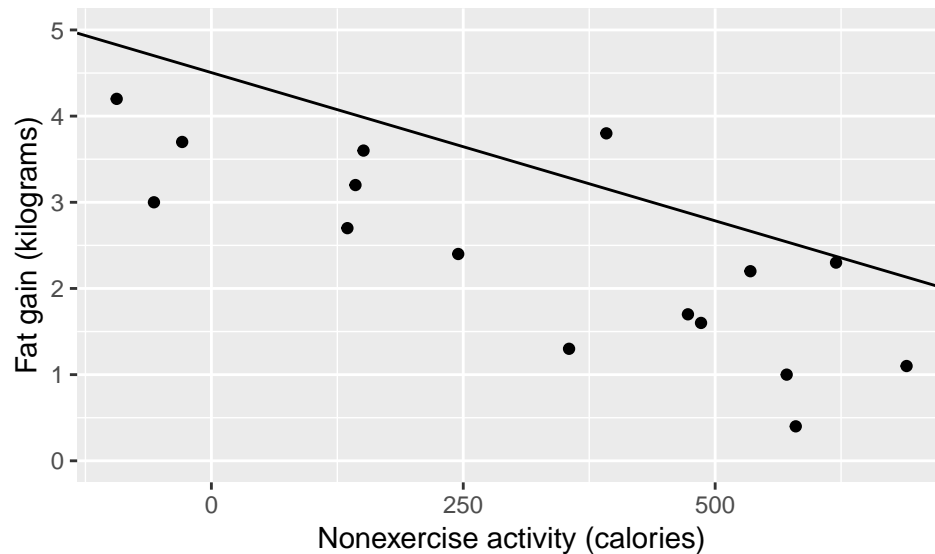
```
gf_point(Fat ~ NEA, data = Fidgeting) %>%
  gf_lm() %>%
  gf_labs(x = "Nonexercise activity (calories)", y = "Fat gain (kilograms)")
```



Use Your Knowledge 2.61: Plot the line

```
gf_point(Fat ~ NEA, data = Fidgeting) %>%
```

```
gf_abline(slope = -.00344, intercept = 4.505) %>%
gf_labs(x = "Nonexercise activity (calories)", y = "Fat gain (kilograms)") +
ylim(0, 5)
```



Example 2.21: Prediction for fat gain

```
fatlm <- lm(Fat ~ NEA, data = Fidgeting)
fatfun <- makeFun(fatlm)
fatfun(NEA = 400)
```

```
##          1
## 2.128528
```

Example 2.24: Regression

```
# page 113
msummary(fatlm)
```

```
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  3.5051229  0.3036164  11.545 1.53e-08 ***
## NEA         -0.0034415  0.0007414  -4.642 0.000381 ***
##
## Residual standard error: 0.7399 on 14 degrees of freedom
## Multiple R-squared:  0.6061, Adjusted R-squared:  0.578
## F-statistic: 21.55 on 1 and 14 DF,  p-value: 0.000381
```

Example 2.25: Fidgeting and fat gain


```
# XX NH Issue with plot and ranges
NEAlm <- lm(NEA ~ Fat, data = Fidgeting) # using fat gain to predict nonexercise activity
NEAfun <- makeFun(NEAlm)
NEAfun(Fat = 6) # (-311.5197, 6)
```

```
##          1
## -311.5197
```

```
NEAfun(Fat = 0) # (745.2604, 0)
```

```
##          1
## 745.2604
```

```
# point 1
# Figure 2.20, page 115
gf_point(Fat ~ NEA, data = Fidgeting) %>%
  gf_lm() %>%
  #gf_fun(NEAfm, color = "red") +
  gf_segment(6 + 0 ~ -311.5197 + 745.2604, color = "red") %>%
  gf_labs(x = "Nonexercise activity (calories)", y = "Fat gain (kilograms)") +
  ylim(0, 6)
```



Section 2.5: Cautions about Correlation and Regression

Example 2.26: Residuals for fat gain

```
fatlm
```

```
##
## Call:
## lm(formula = Fat ~ NEA, data = Fidgeting)
##
## Coefficients:
## (Intercept)          NEA
##    3.505123    -0.003441
```

```
fatfun(NEA = 135)
```


```
##          1
## 3.040522
```

```
2.7 - fatfun(NEA = 135)
```


```
##          1
## -0.3405222
```

Figure 2.23, page 124

```
gf_point(Fat ~ NEA, data = Fidgeting) %>%
  gf_lm() %>%
  gf_labs(x = "Nonexercise activity (calories)", y = "Fat gain (kilograms)")
```



```
gf_point(resid(fatlm) ~ NEA, data = Fidgeting) %>%
  gf_lm() %>%
  gf_labs(x = "Nonexercise activity (calories)", y = "Residual")
```




Example 2.27: Patterns in birthrate and Internet user residuals

```
IntBirth <- read_csv("https://nhorton.people.amherst.edu/ips9/data/chapter02/EG02-027INBIRTH.csv")


## Parsed with column specification:
## cols(
##   Country_Name = col_character(),
##   CountryCode = col_character(),
##   BirthRate2011 = col_double(),
```

```
## UsersPreviousYear = col_double(),
## Users = col_double(),
## LogBirth = col_double(),
## LogUsers = col_double()
## )

intbirthlm <- lm(BirthRate2011 ~ Users, data = IntBirth)
# Figure 2.24, page 126
gf_point(BirthRate2011 ~ Users, data = IntBirth) %>%
  gf_lm() %>%
  gf_labs(x = "Internet users per 100 people", y = "Births per 1000 people")
```



```
gf_point(resid(intbirthlm) ~ Users, data = IntBirth) %>%
  gf_lm() %>%
  gf_labs(x = "Internet users per 100 people", y = "Residual")
```



Example 2.28: Diabetes and blood sugar



```
Diabetes <- read_csv("https://nhorton.people.amherst.edu/ips9/data/chapter02/EG02-028HBA1C.csv")
```

```
## Parsed with column specification:
## cols(
##   Subject = col_integer(),
##   HbA1c_percent = col_double(),
##   FPG_mg_ml = col_integer()
## )
```


```
diabeteslm <- lm(FPG_mg_ml ~ HbA1c_percent, data = Diabetes)
```

```
# Figure 2.25, page 127
```

```
gf_point(FPG_mg_ml ~ HbA1c_percent, data = Diabetes) %>%
  gf_lm() %>%
  gf_labs(x = "HbA1c (percent)", y = "Fasting plasma glucose (mg/dl)")
```



```
gf_point(resid(diabeteslm) ~ HbA1c_percent, data = Diabetes) %>%
  gf_lm() %>%
  gf_labs(x = "HbA1c (percent)", y = "Residual")
```



Example 2.29: Influential observations

Section 2.6: Data Analysis from Two Way Tables

Section 2.7: The Question of Causation