

IPS9 in R: Producing data (Chapter 3)

Shukry Zablah (szablah20@amherst.edu) and Nicholas Horton (nhorton@amherst.edu)

July 18, 2018

Introduction and background

These documents are intended to help describe how to undertake analyses introduced as examples in the Ninth Edition of *Introduction to the Practice of Statistics* (2017) by Moore, McCabe, and Craig.

More information about the book can be found [here](#). The data used in these documents can be found under Data Sets in the Student Site. This file as well as the associated R Markdown reproducible analysis source file used to create it can be found at <https://nhorton.people.amherst.edu/ips9/>.

This work leverages initiatives undertaken by Project MOSAIC (<http://www.mosaic-web.org>), an NSF-funded effort to improve the teaching of statistics, calculus, science and computing in the undergraduate curriculum. In particular, we utilize the `mosaic` package, which was written to simplify the use of R for introductory statistics courses. A short summary of the R needed to teach introductory statistics can be found in the mosaic package vignettes (<http://cran.r-project.org/web/packages/mosaic>). A paper describing the mosaic approach was published in the *R Journal*: <https://journal.r-project.org/archive/2017/RJ-2017-024>.

Chapter 3: Producing data

This file replicates the analyses from Chapter 3: Producing data.

First, load the packages that will be needed for this document:

```
library(mosaic)
library(readr)
```

Section 3.1: Sources of Data

Section 3.2: Design of experiment

See Example 3.12 on page 178.

We are looking to setup a dataset with randomly assigned a “Control” or “Treatment” group to our dataset.

```
set.seed(1) #guarantees the same random results each time R runs
Randomized <- data.frame(ID = c(1:10), randNum = runif(10))
Randomized
```

```
##      ID      randNum
## 1     1 0.26550866
## 2     2 0.37212390
## 3     3 0.57285336
## 4     4 0.90820779
## 5     5 0.20168193
## 6     6 0.89838968
## 7     7 0.94467527
## 8     8 0.66079779
## 9     9 0.62911404
## 10   10 0.06178627
```

We created a data frame in which the first column corresponds to unique IDs and the second column holds the random numbers that will determine which group will be assigned to which ID.

Our plan is to order the IDs based on their random number and assign half of the observations to the Control group and the other half to the Treatment group. We start by arranging the observations.

```
Randomized <- Randomized %>%  
  arrange(randNum)  
Randomized
```

```
##    ID    randNum  
## 1  10 0.06178627  
## 2   5 0.20168193  
## 3   1 0.26550866  
## 4   2 0.37212390  
## 5   3 0.57285336  
## 6   9 0.62911404  
## 7   8 0.66079779  
## 8   6 0.89838968  
## 9   4 0.90820779  
## 10  7 0.94467527
```

And finally we assign each observation to a group.

```
Randomized <- Randomized %>%  
  mutate(Group = c(rep("Treatment", 5), rep("Control", 5)))  
Randomized
```

```
##    ID    randNum    Group  
## 1  10 0.06178627 Treatment  
## 2   5 0.20168193 Treatment  
## 3   1 0.26550866 Treatment  
## 4   2 0.37212390 Treatment  
## 5   3 0.57285336 Treatment  
## 6   9 0.62911404  Control  
## 7   8 0.66079779  Control  
## 8   6 0.89838968  Control  
## 9   4 0.90820779  Control  
## 10  7 0.94467527  Control
```

Now your observations are randomly assigned to a group! Here we chose to assign the first five ID to the Treatment group and the rest to the Control group, but there are other other ways to do it. To finish you can clean up the dataset by selecting only the ID and the Group columns and order it by ID.

```
Randomized %>%  
  select(ID, Group) %>%  
  arrange(ID)
```

```
##    ID    Group  
## 1   1 Treatment  
## 2   2 Treatment  
## 3   3 Treatment  
## 4   4  Control  
## 5   5 Treatment  
## 6   6  Control  
## 7   7  Control  
## 8   8  Control  
## 9   9  Control
```

```
## 10 10 Treatment
```

You can do this for many Groups and for as many observations as you want.

Section 3.3: Sampling Design

We want to get a simple random sample (SRS) of 25 countries out of 189.

```
Countries <- read_csv("https://nhorton.people.amherst.edu/ips9/data/chapter03/EG03-23TTS.csv")
```

In R you just have to use the `sample()` function.

```
Sample <- sample(Countries, size = 25)
Sample
```

```
## # A tibble: 25 x 4
##   CountryName CountryCode   Time orig.id
##   <chr>        <chr>      <int> <chr>
## 1 Costa Rica   CRI          24 39
## 2 Cameroon    CMR          15 34
## 3 Nepal        NPL          17 129
## 4 Croatia      HRV           8 72
## 5 Romania      ROU           9 143
## 6 Korea, Rep. KOR           6 92
## 7 Pakistan     PAK          21 132
## 8 Vietnam      VNM          34 181
## 9 Guyana        GUY          20 69
## 10 Portugal    PRT           3 140
## # ... with 15 more rows
```

Section 3.4: Ethics